

Designing for AI-Powered Social Computing Systems

Gionnieve Lim
gionnievelim@gmail.com
School of Computing, KAIST
Republic of Korea
Singapore University of Technology
and Design
Singapore

Hyunwoo Kim
khw0726@kaist.ac.kr
School of Computing, KAIST
Republic of Korea

Yoonseo Choi
yoonseo.choi@kaist.ac.kr
School of Computing, KAIST
Republic of Korea

Toby Jia-Jun Li
toby.j.li@nd.edu
University of Notre Dame
USA

Chinmay Kulkarni
chinmay.kulkarni@emory.edu
Emory University
USA

Hariharan Subramonyam
harihars@stanford.edu
Stanford University
USA

Joseph Seering
seeringj@stanford.edu
Stanford University
USA

Michael S. Bernstein
msb@cs.stanford.edu
Stanford University
USA

Amy X. Zhang
axz@cs.uw.edu
University of Washington
USA

Elena L. Glassman
glassman@seas.harvard.edu
Harvard University
USA

Simon Perrault
perrault.simon@gmail.com
Singapore University of Technology
and Design
Singapore

Juho Kim
juhokim@kaist.ac.kr
School of Computing, KAIST
Republic of Korea

ABSTRACT

The CSCW community has been active in designing, implementing, and evaluating novel social computing systems. In recent years, there has been a rise in using AI to empower social interactions and the capabilities of these systems. While these implementations charge ahead of the establishment of ethical and legal frameworks, it is timely to reflect on the state of AI-powered social computing systems and to identify new research agendas for the community. This Special Interest Group aims to bring in researchers and practitioners from different fields to foster discussions on the key considerations and challenges in designing for AI-powered social computing systems and to promote opportunities for new research collaborations.

CCS CONCEPTS

• **Human-centered computing** → HCI theory, concepts and models; Collaborative and social computing systems and tools.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
CSCW '23, October 14–18, 2023, Minneapolis, MN, USA

© 2023 Association for Computing Machinery.
ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00
<https://doi.org/XXXXXXXX.XXXXXXX>

KEYWORDS

artificial intelligence, social computing systems, design, human-AI interaction

ACM Reference Format:

Gionnieve Lim, Hyunwoo Kim, Yoonseo Choi, Toby Jia-Jun Li, Chinmay Kulkarni, Hariharan Subramonyam, Joseph Seering, Michael S. Bernstein, Amy X. Zhang, Elena L. Glassman, Simon Perrault, and Juho Kim. 2023. Designing for AI-Powered Social Computing Systems. In *The 26th ACM Conference On Computer-Supported Cooperative Work And Social Computing, October 14–18, 2023, Minneapolis, MN, USA*. ACM, New York, NY, USA, 4 pages. <https://doi.org/XXXXXXXX.XXXXXXX>

1 AI-POWERED SOCIAL COMPUTING SYSTEMS

Social computing systems are online platforms that facilitate interaction and collaboration between people [15]. These systems enable large groups of users to connect with one another, share information and resources, and engage in collaborative activities such as discussing, problem-solving, and decision-making. Some examples of social computing systems include social media platforms, online communities, and collaborative tools. With the increasing capabilities of Artificial Intelligence (AI) and the release of powerful and accessible application programming interfaces like OpenAI's set of GPT models¹, there has been a growing number of applications that use AI to support user interactions. In the field of CSCW research, there is a line of research proposing novel AI-powered social computing systems or integrating AI in existing systems to

¹<https://openai.com/product>

enhance workflows and processes both between humans and with the AI.

For instance, StarryThoughts [6] uses a recommendation algorithm to support the public in exploring and understanding a diverse range of opinions on social issues. Wikum [19] fosters collaboration among readers and editors to construct discussion summaries that are powered by an extractive summarization algorithm. SolutionChat supports moderating real-time structured discussions by recommending appropriate chat messages to the moderators [9]. There are also systems that support the collaborative design process of AI like ModelLens [5] which visualizes model errors and enables developers to configure the system using their own error annotation ontology. These examples from the CSCW research show the diversity of AI technologies and their unlimited potential of applying to social computing systems to improve existing workflows, engage social interactions and enhance collaborations. As AI technologies will only see greater use in the social computing scene, it is timely for the community to take stock of the key considerations, challenges, and opportunities in designing AI-powered social computing systems.

2 ETHICAL CONSIDERATIONS OF USING AI IN SOCIAL COMPUTING SYSTEMS

CSCW and related communities such as FAccT and CHI have been actively researching the problems with applying AI in real-world settings. For example, the lack of transparency makes it challenging to use AI in critical decision-making scenarios where wrong decisions can have devastating consequences [20]. AI also falls short in real-life decision-making where intangible ethical and moral human factors are called for to guide decisions [4].

Introducing AI into social computing systems has created another set of societal challenges. For instance, social media recommendation algorithms meant to increase user engagement have instead led to increasing polarization and ignorance among users [13]. Reflecting the social preferences and prejudices of its designers, AI has also been found to exhibit unfair discrimination in its decision-making towards particular groups, often those that are underrepresented [12]. These problems largely went unnoticed during the design and implementation of the AI but were illuminated only after experimentation and examination by the public post-deployment.

Recent advances in generative AI have also posed new challenges as it is becoming increasingly hard to distinguish between humans and AI and there is even the potential for AI to replace human roles in social interactions. For example, Park et al. introduced 'Social Simulacra' [14], a novel prototyping technique for testing social interactions on social computing systems by simulating the dynamics of social interactions from the given community design, such as the participants' personas and the community rules. In the same vein, SyntheticUsers² was developed to conduct user testing without users.

What are the implications of these developments in social computing systems and to researchers and practitioners who are the key drivers of the field? Can we trust AI to assess or operate the systems, even taking humans out of the equation? Which role should AI serve and how can we ensure that it is achieved?

²<https://www.syntheticusers.com/>

3 TOWARDS UNDERSTANDING AND IMPLEMENTING AI IN SOCIAL COMPUTING SYSTEMS

Researchers have investigated tools and guidelines to enhance users' understanding and designers' consideration of AI in social computing systems. Epstein et al. [3] looked at how presenting explanations increased the effectiveness of human-AI warnings on false content. Lam et al. presented IndieLabel [7], a tool for lay users to audit moderation algorithms by predicting users' perspectives from a small set of user-generated labels of content and using the predicted perspectives to locate potential disagreements between the users and the algorithm. Shen et al. [17] developed Model Card Authoring Toolkit for community members to make a collective decision on adopting an appropriate ML model aligned with community values by supporting the members to be informed and deliberate about the ML models and their tradeoffs. To understand user needs for explainability, Liao et al. [11] developed an algorithm-informed explainable AI (XAI) question bank filled with prototypical questions driven by interviewing UX and design practitioners.

A large body of research has been done to embrace diverse stakeholders' voices to inform the design of AI. Cambre et al. [1] engaged participants in a crowd work platform to speculate about the future of AI voice assistants by completing stories on various everyday settings in 2050, illuminating potential capabilities and concerns with this technology as imagined by the participants. Choi et al. [2] studied creators' perceptions of YouTube algorithms to propose how algorithmic, socioeconomic platforms could be designed to be more creator centered. Seering et al. [16] interviewed moderators across various social media platforms to understand how meaningful communities can be developed, both with and without algorithmic support.

Beyond being informants, another body of work looks at involving stakeholders in the design process of AI-based platforms. WeBuildAI [8] proposed a framework allowing stakeholders in the gig economy to design algorithmic policies and build computational models. Zhang et al. [18] conducted participatory design sessions with drivers to re-design algorithmic management features such as collective driver data sharing. To fight against AI inequality in gig work, Li et al. [10] proposed a bottom-up approach for gig workers about specifying goals, supporting work planning, and sharing data together.

Much has been done to understand the impacts of AI on users and there is a move towards being more mindful in the application of AI. Yet, with social dynamics at play in social computing systems, what further considerations are needed when incorporating AI? To what extent should AI be used? How can the effects of AI be measured and managed?

4 SIG GOALS

Developments of AI have reached a new level of maturity with recent releases of publicly available AI technologies like OpenAI's GPT models and MidJourney which has produced new potentials and also problems for using AI in social computing systems. As such, we see CSCW'23 as an integral milestone for the community to reflect on and contribute their diverse perspectives, expertise, and

opinions on the considerations and advancements of AI-powered social computing systems.

This SIG invites sharing and discussions on designing AI-powered social computing systems to provide a space for the CSCW community to reflect on the challenges and opportunities ahead. We aim to connect participants with diverse backgrounds and perspectives to promote a holistic discussion on the societal and technological issues of AI-powered social computing systems. We aim to build a community of interested members and to generate a set of research agendas that will be made publicly available afterward.

5 THEMES

In social computing systems, the inclusion of AI adds another layer of complexity to the social dynamics already at play. In social media, AI can foster engagement yet also amplify undesired societal effects. In collaborative systems, AI can affect team dynamics and decisions. The power of AI in influencing social dynamics for better or worse poses a significant challenge for researchers and practitioners, one that calls for scrutiny when considering the design choices of the AI that is used in social computing systems.

To support and expand the discussions on designing AI-powered social computing systems, the following are themes of the SIG:

- What are the key considerations and challenges in designing AI-powered social computing systems involving stakeholders?
- To what extent should AI be used in social computing systems and what are the intended (or potentially unintended) outcomes?
- What methods are appropriate to assess and address these outcomes?
- How has the advent of generative AI models changed the landscape of social computing systems and what opportunities are there?

6 INTENDED COMMUNITY

The SIG will be relevant to researchers and practitioners of the CSCW community who have an interest in:

- building novel AI-powered social computing systems
- auditing the uses of AI in social computing systems
- examining changes in the social computing landscape with new technological developments

While these members may have greater interest, the event will be open to all attendees. The SIG welcomes both experienced members of the community wishing to develop the field further and newcomers seeking to understand the field. Social computing systems permeate everyday life and with AI coming to the fore, we hope that the discussions we facilitate will be interesting and meaningful to the broader audience.

7 FORMAT

Activities of the SIG will begin around a month before the event. A Slack workspace will be created to connect attendees and serve as a common space for further discussions after the SIG.

7.1 Before the SIG

We will send an open call through various social media channels of organizers to spread the word. Along with the call, we will embed a link to a survey to solicit open questions on the themes of the SIG and suggestions for its format. The questions will be used to seed discussions during the SIG while we will consider the suggestions to improve the format of the SIG. An invite to the Slack workspace will also be embedded and we will encourage attendees to introduce themselves and to share a relevant paper outside of the CSCW field that they find important to designing for AI in social computing systems. These calls will be sent several times in the period leading up to the SIG for the greatest visibility. In addition, we will send information about the SIG to colleagues who would have an interest in the topic and encourage them to further share the call.

7.2 During the SIG

The SIG will be held in-person with a duration of 75 minutes. To ensure a focused and fruitful discussion, we propose the following plan:

- Introduction (5 minutes): We will start with a brief introduction of the co-organizers and their work related to AI-powered social computing systems and provide an overview of the topic, format, and goals of the SIG. Attendees will introduce themselves via a shared document.
- Breakout Session (15 minutes): We will then conduct small group discussions based on the themes of the SIG and the list of questions we collected from the pre-SIG activity.
- Presentation Preparation (10 minutes): After the discussions, each group will choose a final theme to present and further discuss while working on a shared document to prepare their presentation slide.
- Presentation and Q&A (25 minutes): Each group will have 3 minutes to present their chosen theme and the main points of their discussions with a 2-minute Q&A from the audience. The organizers will note down items that have been discussed several times across the groups.
- Floor Discussion (15 minutes): We will reconvene and discuss the issues that have surfaced and the potential solutions for them.
- Closing (5 minutes): We will close with a summary of the discussions and announce that the materials will be made publicly available. Attendees will be invited to join the Slack workspace to continue their networking and discussions.

7.3 After the SIG

We will curate the discussions and materials from the SIG and share them through Slack and via social media channels two weeks after the event. We hope that the document serves to rekindle reflections on the discussions that occurred during the SIG and that it inspires new research directions and further investigations from the community. There is also the possibility of extending the SIG into a journal article or workshop. The Slack workspace will be kept open as a place for sharing relevant publications and for supporting communications among researchers and practitioners with similar interests.

8 LOGISTICS

Participants are encouraged to bring their laptops as there will be online activities. We will also prepare papers and markers for those who do not have access to their devices. The papers will be collected at the end of the SIG and digitized for sharing. The contents will be manually transcribed from handwriting to text and only the text version will be shared.

REFERENCES

- [1] Julia Cambre, Samantha Reig, Queenie Kravitz, and Chinmay Kulkarni. 2020. "All Rise for the AI Director": Eliciting Possible Futures of Voice Technology through Story Completion. In *Proceedings of the 2020 ACM Designing Interactive Systems Conference* (Eindhoven, Netherlands) (*DIS '20*). Association for Computing Machinery, New York, NY, USA, 2051–2064. <https://doi.org/10.1145/3357236.3395479>
- [2] Yoonseo Choi, Eun Jeong Kang, Min Kyung Lee, and Juho Kim. 2023. Creator-Friendly Algorithms: Behaviors, Challenges, and Design Opportunities in Algorithmic Platforms. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems* (Hamburg, Germany) (*CHI '23*). Association for Computing Machinery, New York, NY, USA, Article 564, 22 pages. <https://doi.org/10.1145/3544548.3581386>
- [3] Ziv Epstein, Nicolo Foppiani, Sophie Hilgard, Sanjana Sharma, Elena Glassman, and David Rand. 2022. Do Explanations Increase the Effectiveness of AI-Crowd Generated Fake News Warnings? *Proceedings of the International AAAI Conference on Web and Social Media* 16, 1 (May 2022), 183–193. <https://doi.org/10.1609/icwsm.v16i1.19283>
- [4] Trystan S. Goetze. 2022. Mind the Gap: Autonomous Systems, the Responsibility Gap, and Moral Entanglement. In *2022 ACM Conference on Fairness, Accountability, and Transparency* (Seoul, Republic of Korea) (*FAccT '22*). Association for Computing Machinery, New York, NY, USA, 390–400. <https://doi.org/10.1145/3531146.3533106>
- [5] Yannis Katsis and Christine T. Wolf. 2019. ModelLens: An Interactive System to Support the Model Improvement Practices of Data Science Teams. In *Conference Companion Publication of the 2019 on Computer Supported Cooperative Work and Social Computing* (Austin, TX, USA) (*CSCW '19*). Association for Computing Machinery, New York, NY, USA, 9–13. <https://doi.org/10.1145/3311957.3359512>
- [6] Hyunwoo Kim, Haesoo Kim, Kyung Je Jo, and Juho Kim. 2021. StarryThoughts: Facilitating Diverse Opinion Exploration on Social Issues. *Proc. ACM Hum.-Comput. Interact.* 5, CSCW1, Article 66 (apr 2021), 29 pages. <https://doi.org/10.1145/3449140>
- [7] Michelle S Lam, Mitchell L Gordon, Danaë Metaxa, Jeffrey T Hancock, James A Landay, and Michael S Bernstein. 2022. End-User Audits: A System Empowering Communities to Lead Large-Scale Investigations of Harmful Algorithmic Behavior. *Proceedings of the ACM on Human-Computer Interaction* 6, CSCW2 (2022), 1–34.
- [8] Min Kyung Lee, Daniel Kusbit, Anson Kahng, Ji Tae Kim, Xinran Yuan, Allissa Chan, Daniel See, Ritesh Noothigattu, Siheon Lee, Alexandros Psomas, et al. 2019. WeBuildAI: Participatory framework for algorithmic governance. *Proceedings of the ACM on Human-Computer Interaction* 3, CSCW (2019), 1–35.
- [9] Sung-Chul Lee, Jaeyoon Song, Eun-Young Ko, Seongho Park, Jihee Kim, and Juho Kim. 2020. Solutionchat: Real-time moderator support for chat-based structured discussion. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*. 1–12.
- [10] Toby Jia-Jun Li, Yuwen Lu, Jaylexia Clark, Meng Chen, Victor Cox, Meng Jiang, Yang Yang, Tamara Kay, Danielle Wood, and Jay Brockman. 2022. A Bottom-Up End-User Intelligent Assistant Approach to Empower Gig Workers against AI Inequality. In *2022 Symposium on Human-Computer Interaction for Work* (Durham, NH, USA) (*CHIWORK 2022*). Association for Computing Machinery, New York, NY, USA, Article 9, 10 pages. <https://doi.org/10.1145/3533406.3533418>
- [11] Q. Vera Liao, Daniel Gruen, and Sarah Miller. 2020. Questioning the AI: Informing Design Practices for Explainable AI User Experiences. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) (*CHI '20*). Association for Computing Machinery, New York, NY, USA, 1–15. <https://doi.org/10.1145/3313831.3376590>
- [12] Eirini Ntoutsi, Pavlos Fafalios, Ujwal Gadiraju, Vasileios Iosifidis, Wolfgang Nejdl, Maria-Esther Vidal, Salvatore Ruggieri, Franco Turini, Symeon Papadopoulos, Emmanouil Krasanakis, Ioannis Kompatsiaris, Katharina Kinder-Kurlanda, Claudia Wagner, Fariba Karimi, Miriam Fernandez, Harith Alani, Bettina Berendt, Tina Kruegel, Christian Heinze, Klaus Broelemann, Gjergji Kasneci, Thanassis Tiropanis, and Steffen Staab. 2020. Bias in data-driven artificial intelligence systems—An introductory survey. *WIREs Data Mining and Knowledge Discovery* 10, 3 (2020), e1356. <https://doi.org/10.1002/widm.1356> arXiv:<https://wires.onlinelibrary.wiley.com/doi/pdf/10.1002/widm.1356>
- [13] E. Pariser. 2012. *The Filter Bubble: How the New Personalized Web Is Changing What We Read and How We Think*. Penguin Publishing Group. <https://books.google.co.kr/books?id=2JZPEAAAQBAJ>
- [14] Joon Sung Park, Lindsay Popowski, Carrie Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. 2022. Social Simulacra: Creating Populated Prototypes for Social Computing Systems. In *Proceedings of the 35th Annual ACM Symposium on User Interface Software and Technology* (Bend, OR, USA) (*UIST '22*). Association for Computing Machinery, New York, NY, USA, Article 74, 18 pages. <https://doi.org/10.1145/3526113.3545616>
- [15] Doug Schuler. 1994. Social Computing. *Commun. ACM* 37, 1 (jan 1994), 28–29. <https://doi.org/10.1145/175222.175223>
- [16] Joseph Seering, Tony Wang, Jina Yoon, and Geoff Kaufman. 2019. Moderator engagement and community development in the age of algorithms. *New Media & Society* 21, 7 (2019), 1417–1443. <https://doi.org/10.1177/1461444818821316> arXiv:<https://doi.org/10.1177/1461444818821316>
- [17] Hong Shen, Leijie Wang, Wesley H. Deng, Ciell Brusse, Ronald Velgersdijk, and Haiyi Zhu. 2022. The Model Card Authoring Toolkit: Toward Community-Centered, Deliberation-Driven AI Design. In *2022 ACM Conference on Fairness, Accountability, and Transparency* (Seoul, Republic of Korea) (*FAccT '22*). Association for Computing Machinery, New York, NY, USA, 440–451. <https://doi.org/10.1145/3531146.3533110>
- [18] Angie Zhang, Alexander Boltz, Chun Wei Wang, and Min Kyung Lee. 2022. Algorithmic management reimaged for workers and by workers: Centering worker well-being in gig work. In *CHI Conference on Human Factors in Computing Systems*. 1–20.
- [19] Amy X. Zhang, Lea Verou, and David Karger. 2017. Wikum: Bridging Discussion Forums and Wikis Using Recursive Summarization. In *Proceedings of the 2017 ACM Conference on Computer Supported Cooperative Work and Social Computing* (Portland, Oregon, USA) (*CSCW '17*). Association for Computing Machinery, New York, NY, USA, 2082–2096. <https://doi.org/10.1145/2998181.2998235>
- [20] John Zoshak and Kristin Dew. 2021. Beyond Kant and Bentham: How Ethical Theories Are Being Used in Artificial Moral Agents. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems* (Yokohama, Japan) (*CHI '21*). Association for Computing Machinery, New York, NY, USA, Article 590, 15 pages. <https://doi.org/10.1145/3411764.3445102>